CS 487/587 Database Implementation
Winter 2021
Database Benchmarking Project
Part I - Data Generation & System Selection

Due: Tues February 2, midnight (end of the day on the 2nd), D2L

Team: Swetha Venkatesan, Sai Deepika Gooty

Project Options:

Option 2: Single System: Evaluate a single 'relational' system by changing system parameters and varying relation (data) size.

System Selection: PostgreSQL

Project Part I: Data Generation & System Selection

GitHub Repository - https://github.com/swet09/DB-Implementation-Project

1. Screenshots or other mechanisms that demonstrate that you have successfully generated data and loaded that data into your system.

   Creation of VM

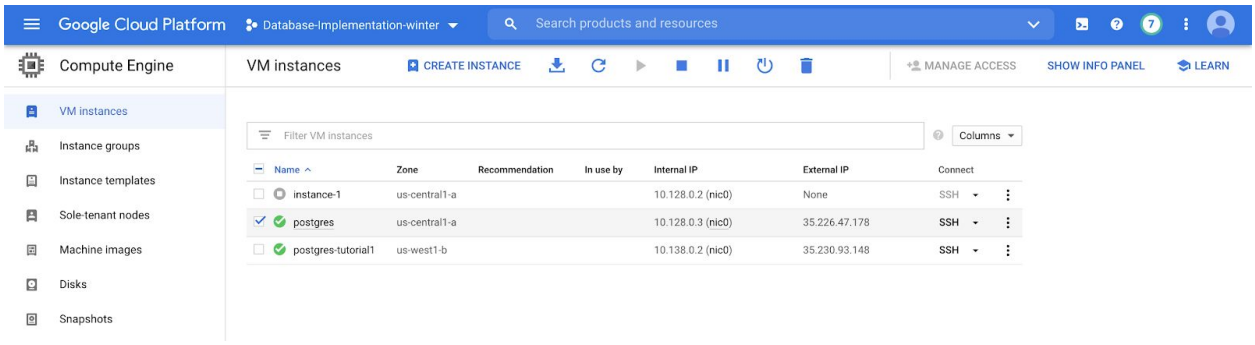

Installation of Postgres

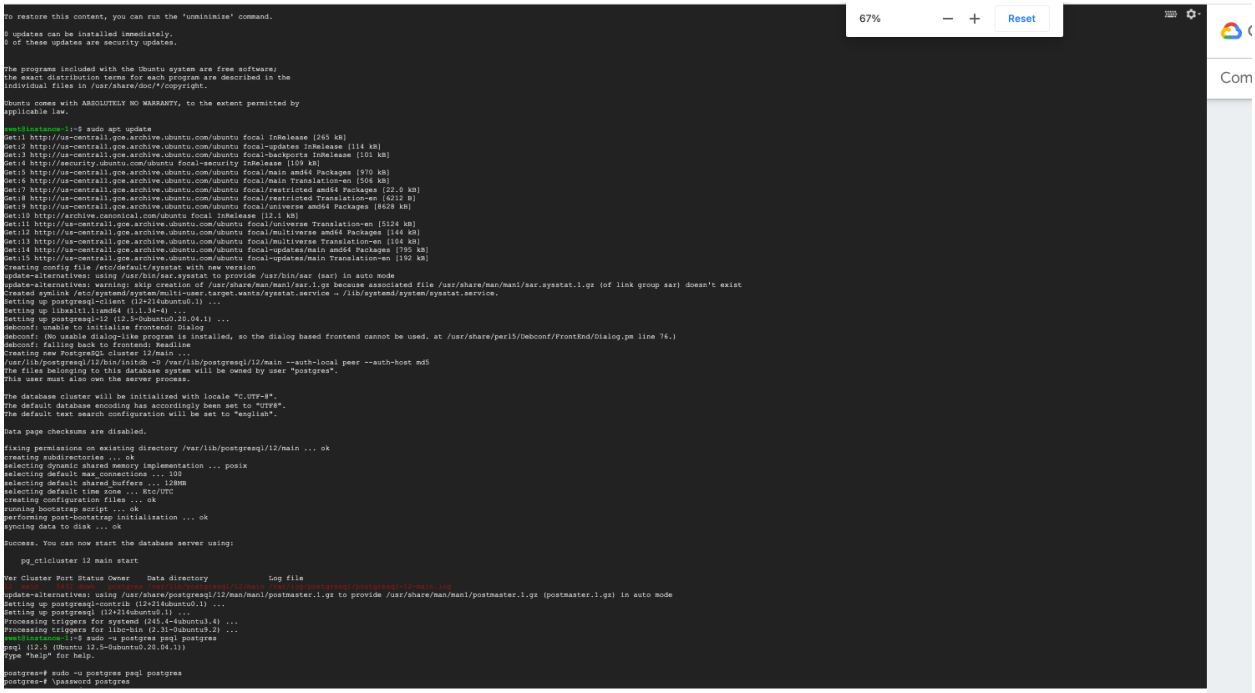

Running the Java Code

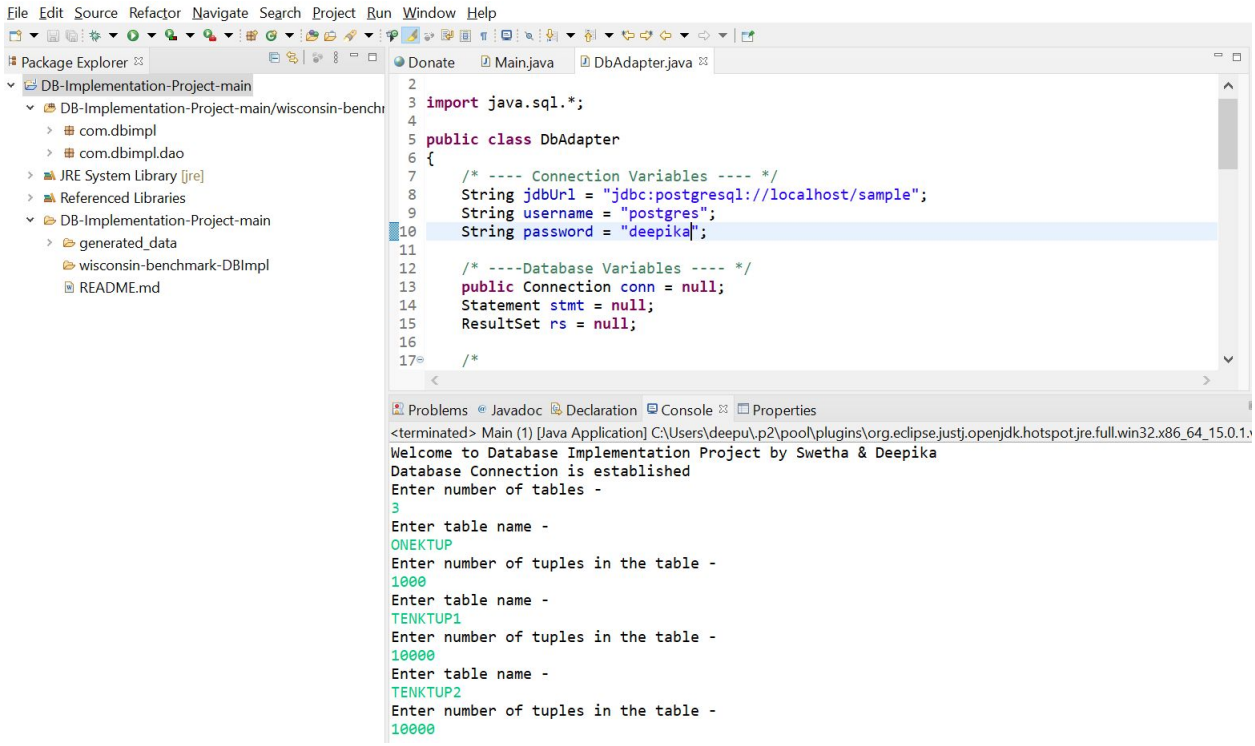

ONEKTUP

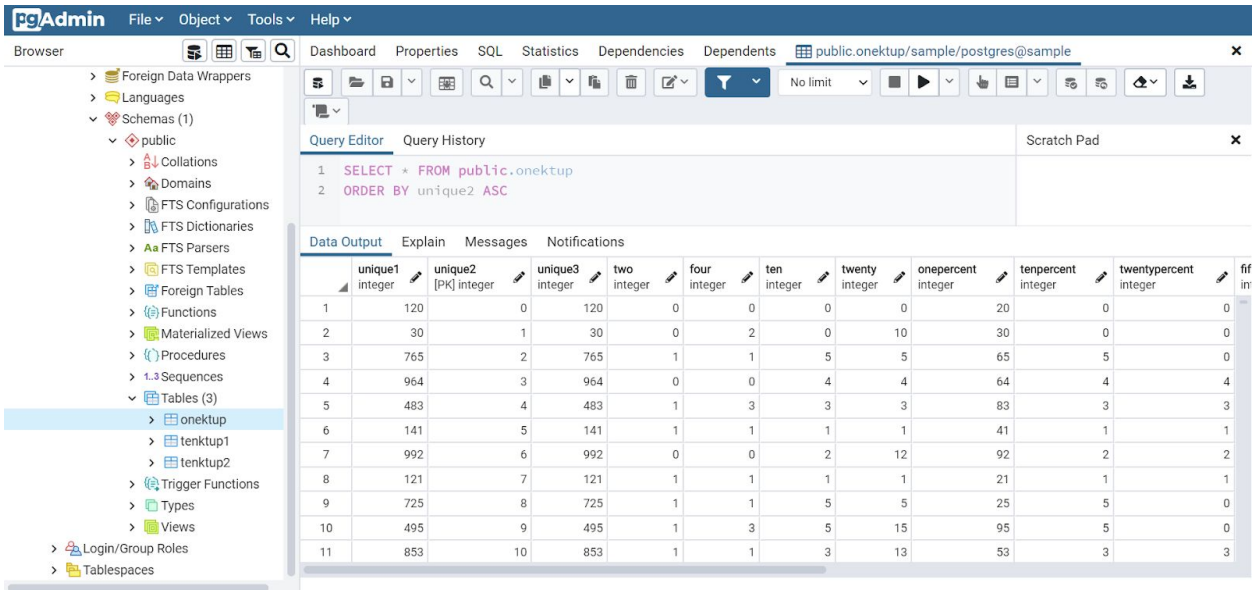

TENKTUP1

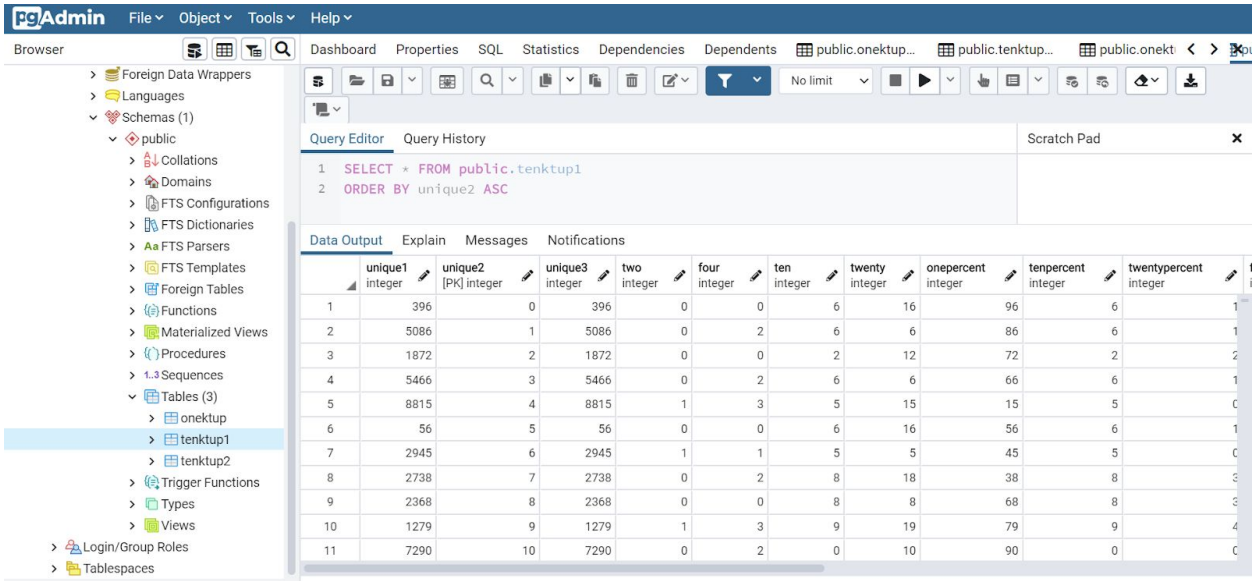

TENKTUP2

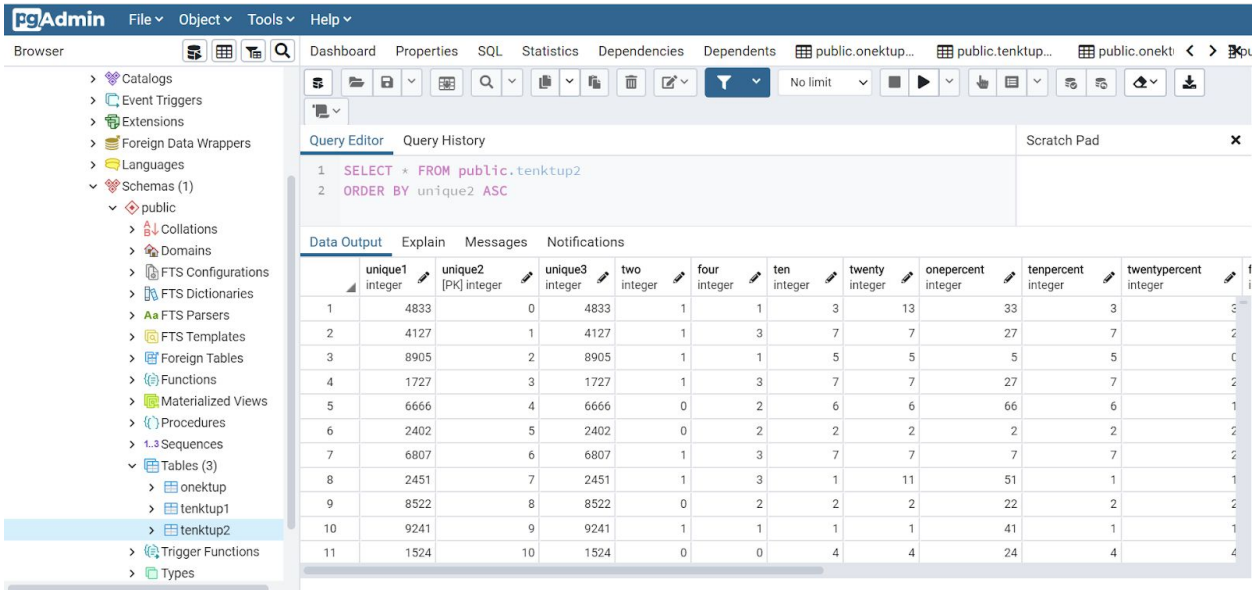

2. README-Part I. 1-2 pages including:
   Link - https://github.com/swet09/DB-Implementation-Project#db-implementation-project

DB-Implementation-Project

Introduction

The project's aim is to evaluate a single relational system such as PostgreSQL by changing system parameters and varying relation or data size and through this project we have an increased understanding of database performance and improve programming experience with a RDBMS system such as PostgreSQL.

Brief description of your data generation process

The main objective of part 1 is to generate data based on the Wiscosin Benchmark specification. The original benchmark was composed of three relations (onektup, tenktup1 and tenktup2), each relation was composed of the thirteen integer attributes and three 52 byte string attributes. The three relations or data are generated using a Java. The 'Main.java' file consists of the main code for data generation and the 'DbAdapter.java' file in the dao folder consists of the Database connection configurations. The generated data files are placed in a folder named ‘generated_data’.

System we are working with and why

The database system we have chosen to work with is PostgreSQL and for data generation we have decided to write code in the programming language Java, because of our prior knowledge and experience with both PostgreSQL and Java is why we decided to go with them. We plan to develop benchmark comparisons using queries

from Wisconsin Benchmark and evaluate against a progress system. We have decided to use a single database system as we want to keep the other parameters such as main memory, OS load and processor speed constant while testing the database against different parameter values and optimizer options.

Prerequisites

- Create a VM on GCP
- Configuring postgres(https://cloud.google.com/community/tutorials/setting-up-postgres)
- Allow remote ip to connect to postgres server

Validation that data is loaded into the system

After successful execution of the java code, data is generated as CSV files in "‘generated_data’ folder as a result of select queries from Postgres.

Lessons learned or issues encountered

The major learnings for this project part 1 was from the Wisconsin Benchmark paper. The paper provided great insight into designing a specification. The information provided by the Benchmark Relation Generator was very useful in generating data files. The important takeaways from the paper were, how the queries in the benchmark should test the performance of the major components of a relational database system and why the semantics and statistics of the underlying relations should be well understood so that it is easy to add new queries and to their behavior. Through this project we learned to generate relations in pgadmin and populate them through java. We did encounter a couple of issues, one was to allow remote ip to connect to postgres server. Another was to randomise the values of unique1. For which we implemented ArrayList and used Collections.shuffle to shuffle that ArrayList. Another issue we faced was while generating values for string4, we used switch case statement to implement it.